Data Flow

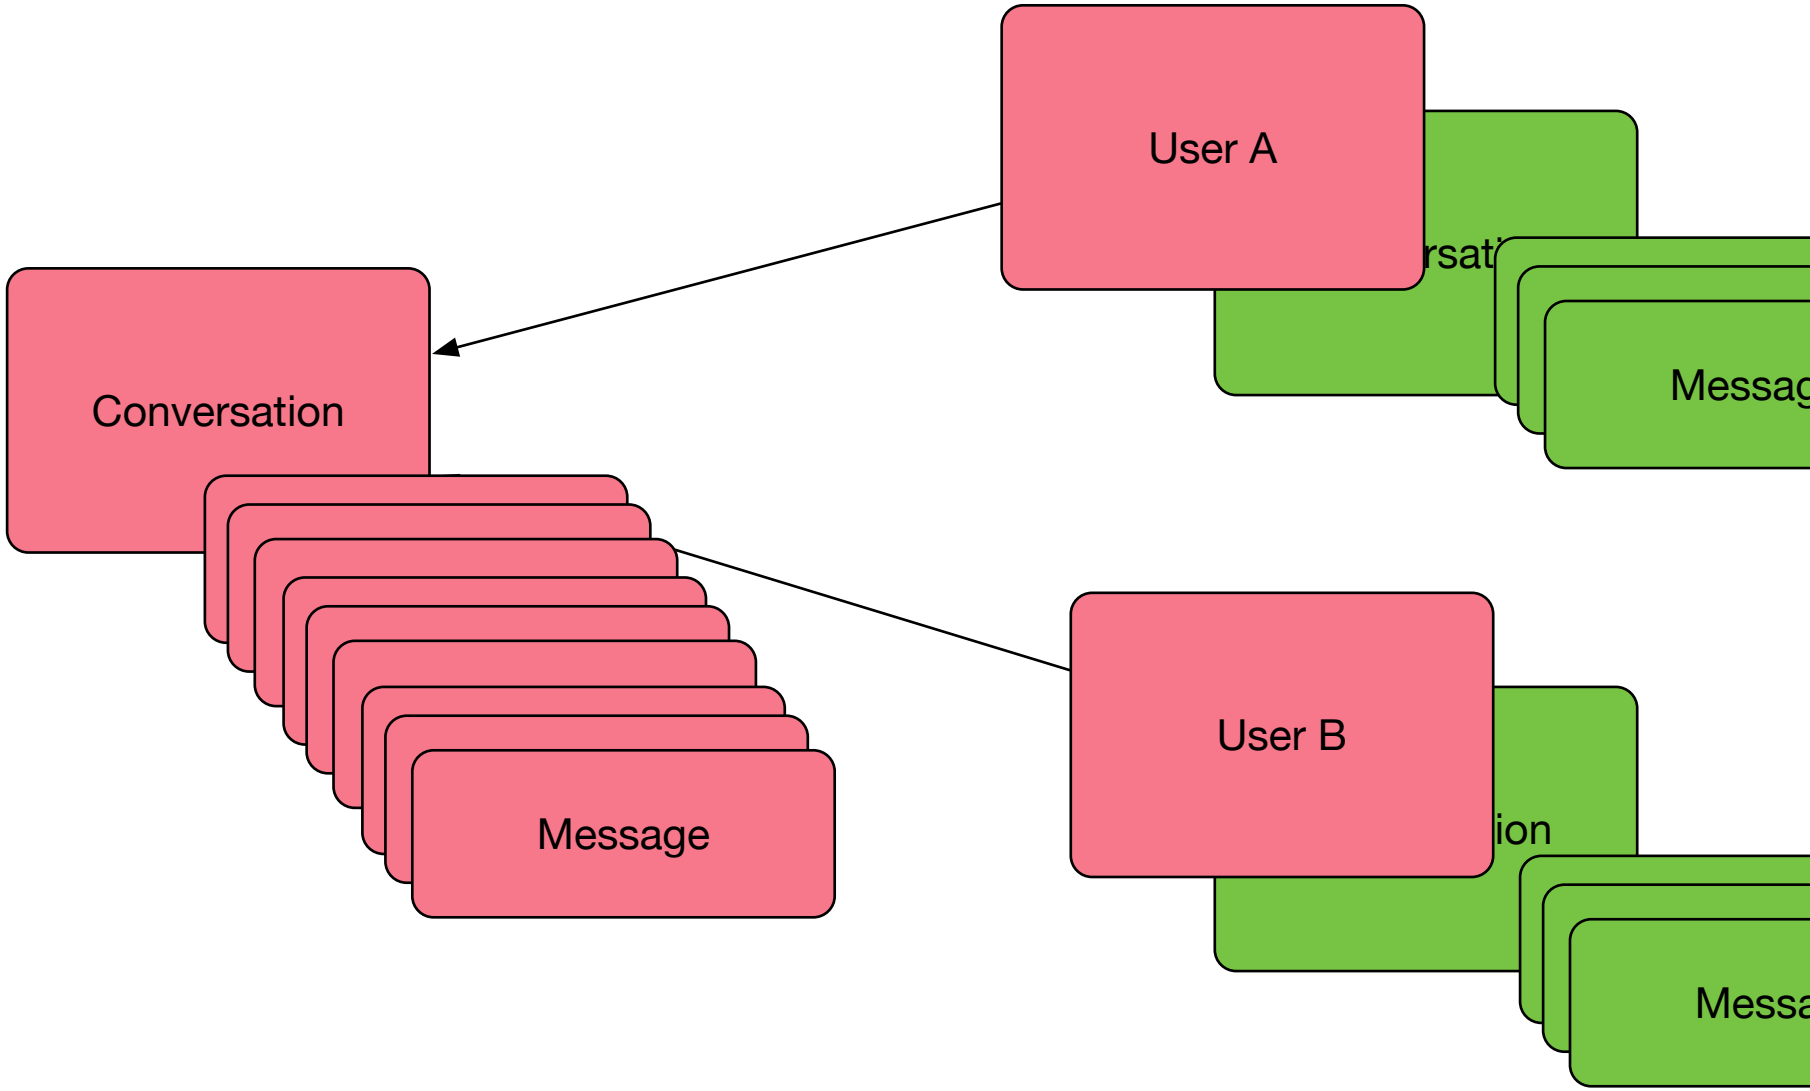

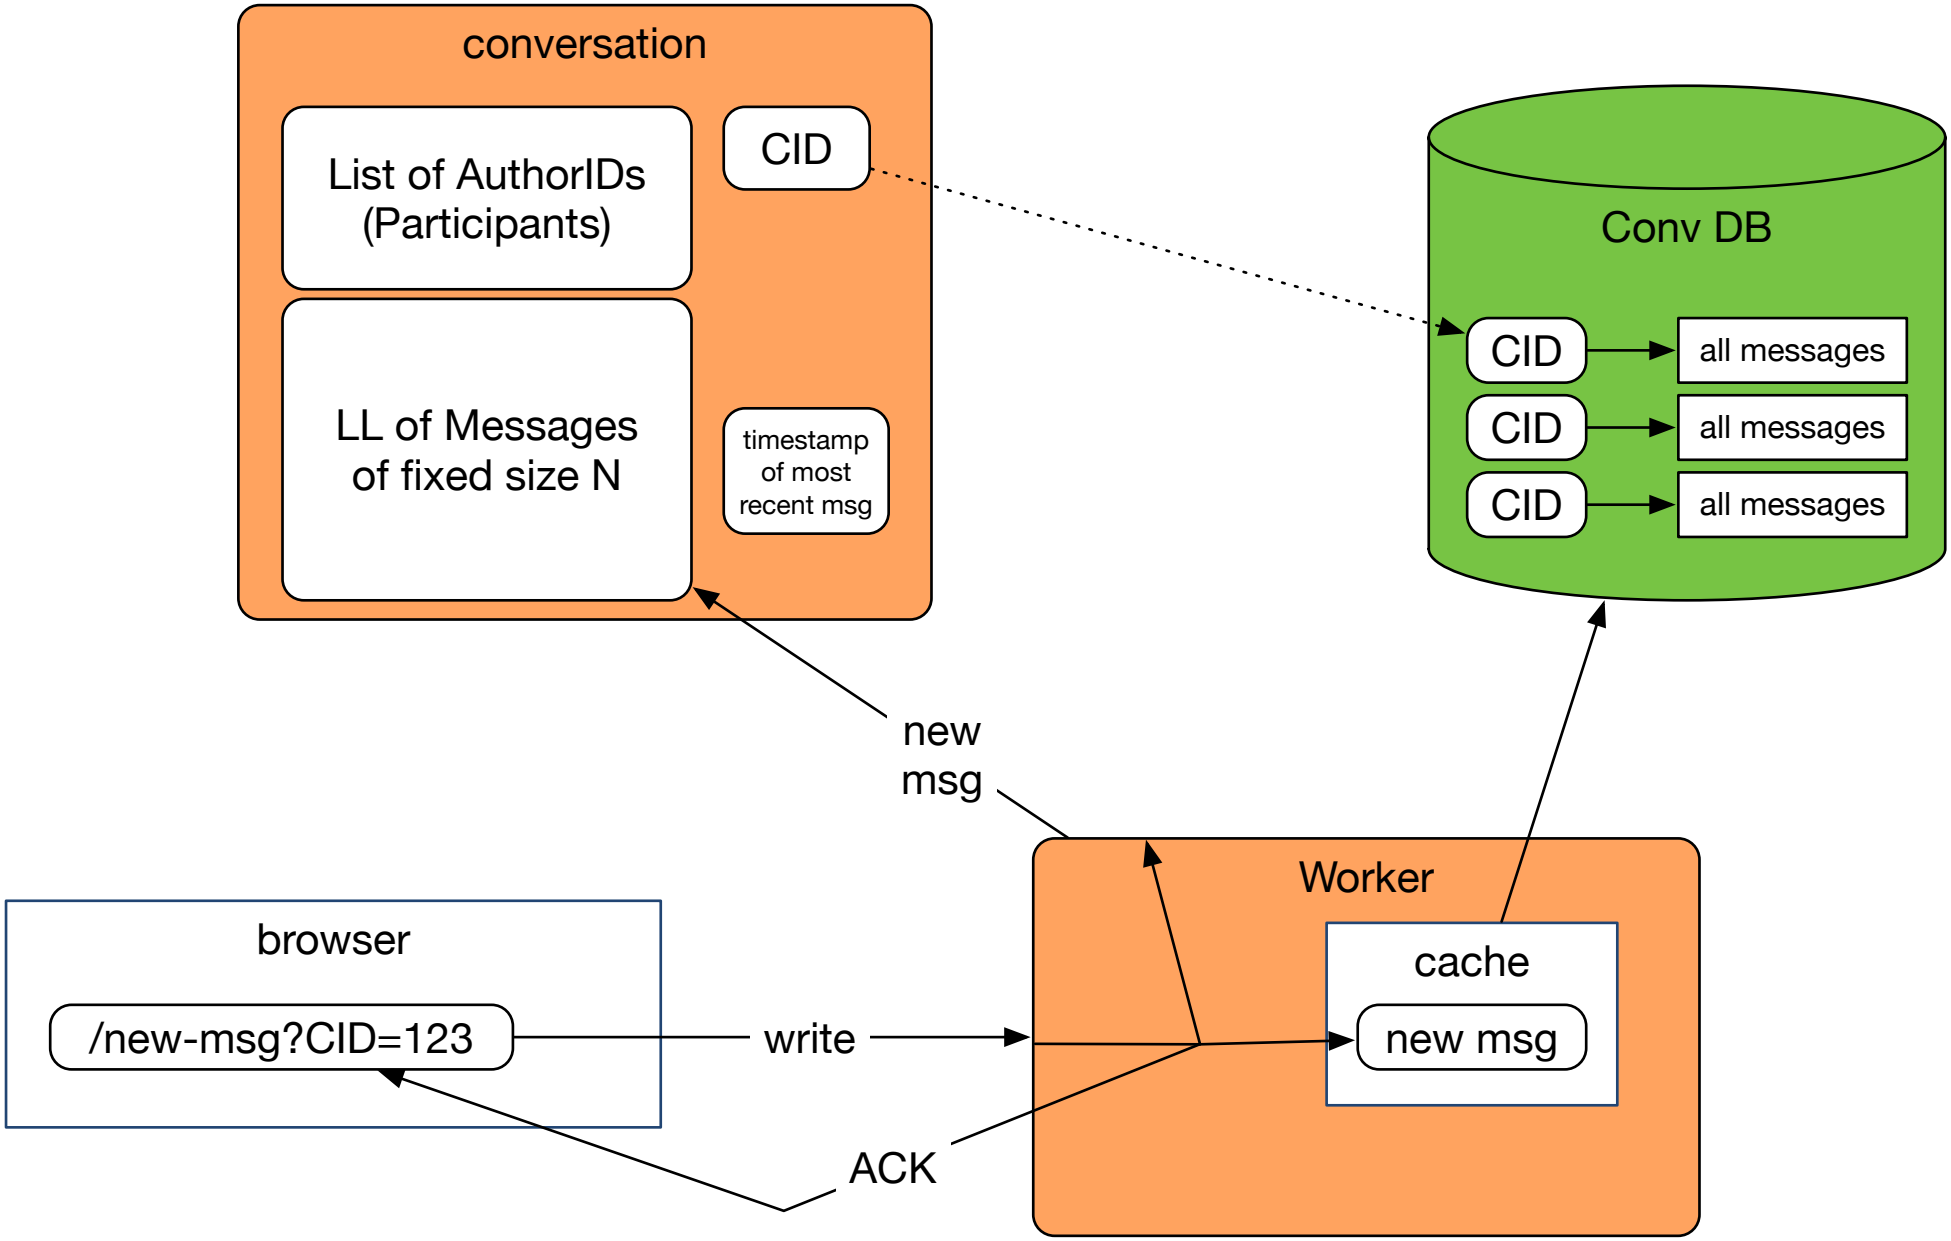

Availability vs “Consistency”

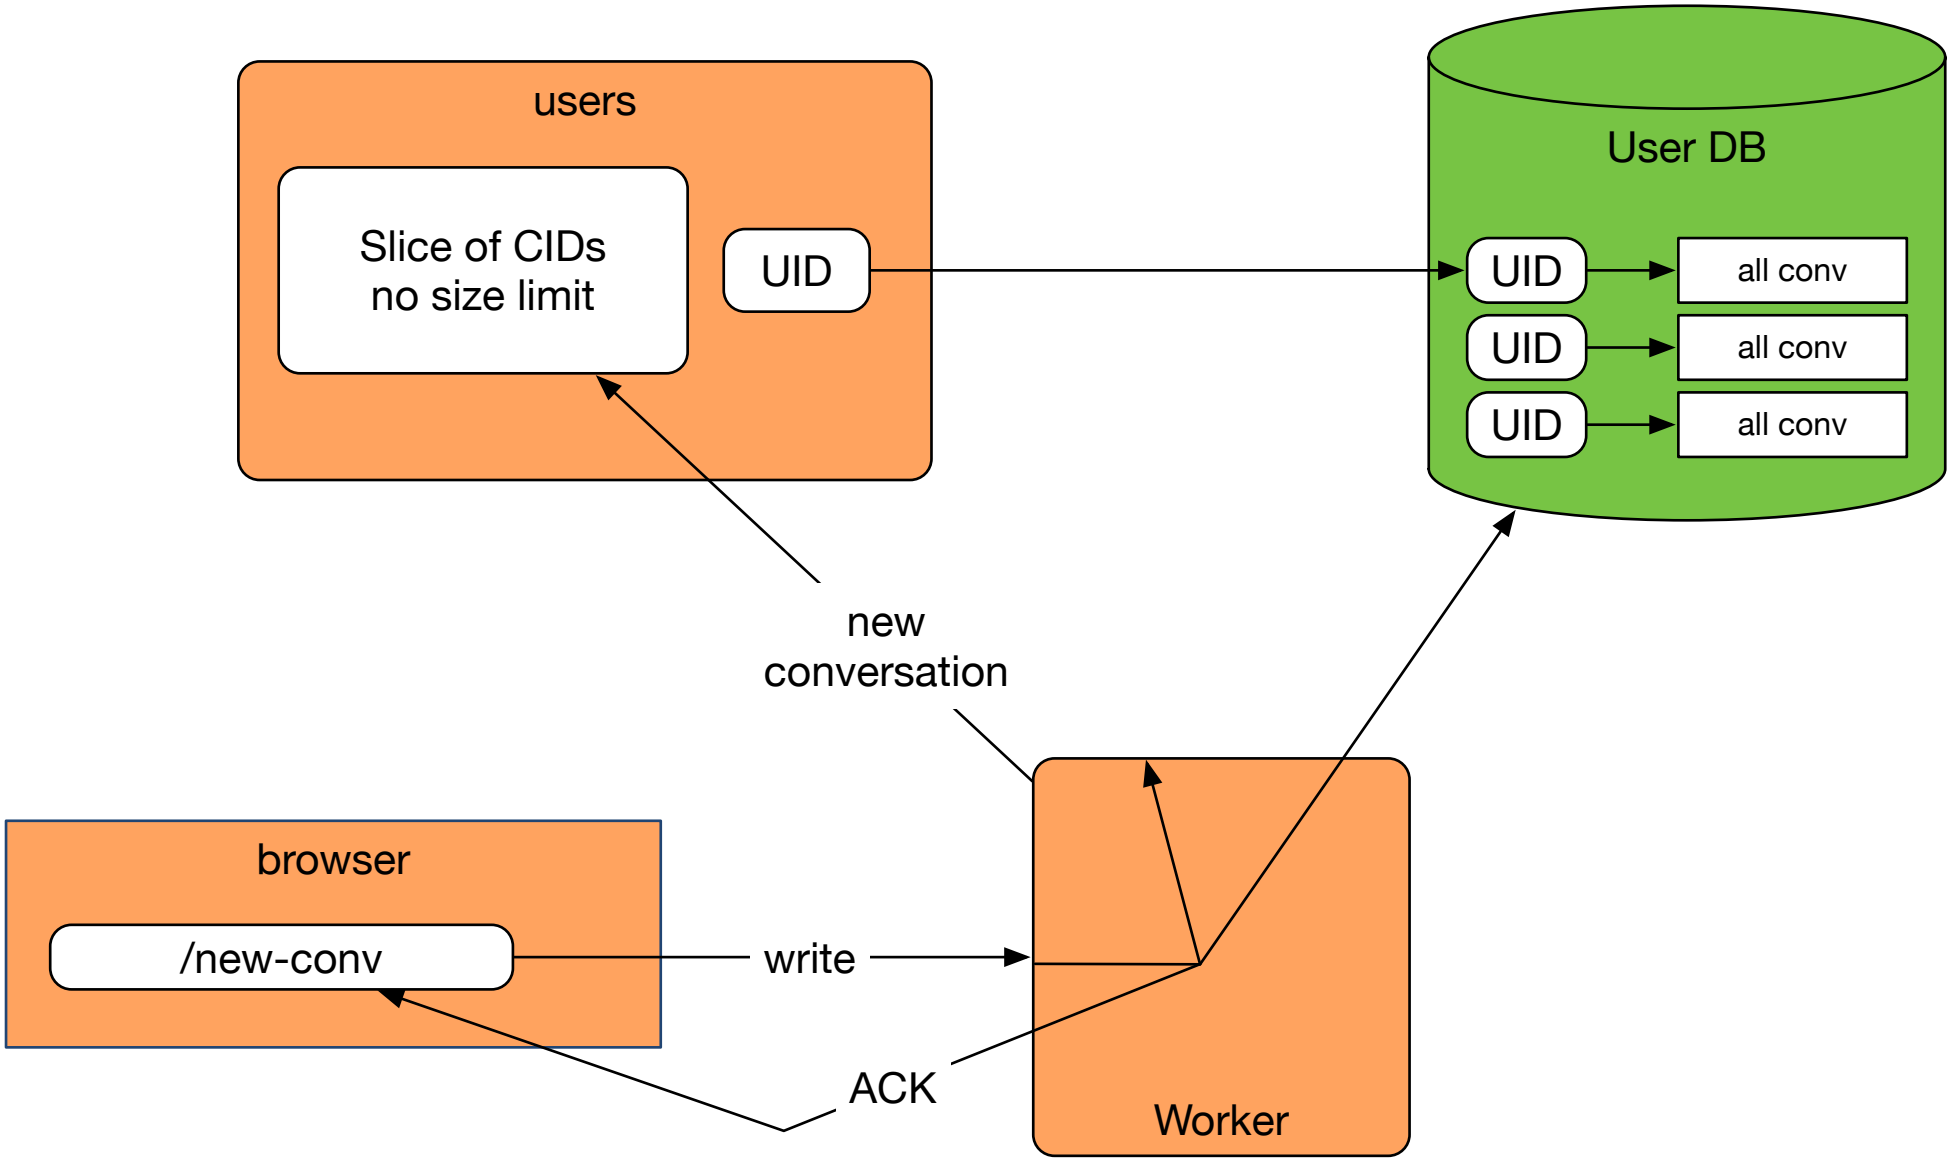
users
Slice of CIDs
no size limit
UID
User DB
UID
all conv
UID
all conv
UID
all conv
new
conversation
browser
/new-conv
write
ACK
Worker

Manages conversations.

Data structures?

```
Map[CID] -> Conversation{}
```

What RPCs can it manage?

“Remote Procedure Call”
*** gRPC —> Google ~10ish years
CORBA / COBRA —> ??? Sun/Oracle?
Thrift —> Facebook ~4ish years
REST —> JSON in response to “raw” http requests
/give/me/cookies [numCookies=4]

```
getConversations(CIDs) // always returns sorted
getConversationsView(CIDs) // always returns sorted
newConversation(UIDs []int32)
newMessage(content string, author int32, CID int32)
deleteConversation(CID int32)
deleteAllTracesOf(UID int32)
```

Manages users.

Data structures?

```
Map[UID] -> User{}
```

What RPCs can it manage?

```
getView(UID int32)
newUser(username string)
deleteUser(UID int32)
```